

InitialNotebook

December 8, 2025

```
[181]: # Mount Google Drive  
from google.colab import drive  
drive.mount('/content/drive/')
```

Drive already mounted at /content/drive/; to attempt to forcibly remount, call
drive.mount("/content/drive/", force_remount=True).

```
[182]: import numpy as np
        import pandas as pd
        import matplotlib.pyplot as plt
        import seaborn as sns
        from sklearn.linear_model import LinearRegression, LogisticRegression
        import statsmodels.api as sm
        from sklearn.preprocessing import StandardScaler
        from scipy.stats import ttest_ind
        from sklearn.feature_selection import SequentialFeatureSelector
        from sklearn.model_selection import train_test_split
```

```
[183]: data = pd.read_csv('/content/drive/MyDrive/Final_Project/  
    ↪shopping_behavior_updated.csv')  
data[0:1]  
#Looking at the data and the columns
```

```
[183]:   Customer ID  Age  Gender  Item Purchased  Category  Purchase Amount (USD) \
          0             1    55    Male      Blouse  Clothing           53

          Location  Size  Color  Season  Review  Rating  Subscription Status  \
          0  Kentucky     L   Gray  Winter     3.1       3.1        Yes

          Shipping Type Discount Applied Promo Code Used  Previous Purchases  \
          0      Express        Yes        Yes        Yes         14

          Payment Method Frequency of Purchases
          0            Venmo        Fortnightly
```

The data has a lot of columns for us to look at. I will begin by playing with different graphs, and looking for any relationships that are important.


Our two goals are to look at what factors increase the **amount of money** spent in a purchase,

and what factors tend to lead to more **customers returning more frequently**.


The factors I initially look at are factors that I suspect will be influential, based on my previous research and experience. This will help me get an idea of how some of the data is spread out, but we will have a better idea when we try the **greedy forward algorithm** to determine which factors are the most influential!

Exploring Data:

```
[184]: plt.figure(figsize=(4, 3))
sns.histplot(data['Gender'])
plt.title('Gender vs. Amount of Purchases')
plt.tight_layout()
plt.show()
```




```
[185]: #lets look at a few columns
#starting with gender and purchase amts.
plt.figure(figsize=(4, 3))
sns.boxplot(x='Gender', y='Purchase Amount (USD)', data=data)
plt.title('Gender vs. Purchase Amount')
plt.show()
```




As seen above, gender does not seem to be a major indicator on the purchase size. We will roughly look at a few other columns before trying other methods to figure out how we can identify which are the most indicative columns.

```
[186]: plt.figure(figsize=(4, 3))
sns.boxplot(x = data['Discount Applied'], y = data['Purchase Amount (USD)'])
plt.title('Discount Applied vs. Purchase Amount')
```


```
[186]: Text(0.5, 1.0, 'Discount Applied vs. Purchase Amount')
```



```
[187]: #using a sample to better see the scatterplot relationships
plt.figure(figsize=(4, 3))
data_sample = data.sample(frac = .5, random_state = 1)
sns.scatterplot(x = data_sample['Age'], y = data_sample['Purchase Amount\u2192(USD)'])
plt.title('Age vs Purchase Amount')
plt.tight_layout()
```




```
[188]: plt.figure(figsize=(4, 3))
sns.boxplot(x = data_sample['Category'], y = data_sample['Purchase Amount\u2192(USD)'])
plt.title('Category vs Purchase Amount')
plt.tight_layout()
```



Of the previous 3 factors, the **category** of the product seems to have more of an influence on the purchase size.


```
[189]: plt.figure(figsize=(4, 3))
sns.histplot(data['Frequency of Purchases'])
plt.xticks(rotation=45)
plt.title('Frequency of Purchases')
```

```
[189]: Text(0.5, 1.0, 'Frequency of Purchases')
```




```
[190]: plt.figure(figsize=(4, 3))
sns.boxplot(x = data['Frequency of Purchases'], y = data['Purchase Amount ↴(USD)'])
plt.xticks(rotation=45)
plt.title('Frequency of Purchases vs Purchase Amount')
```

```
[190]: Text(0.5, 1.0, 'Frequency of Purchases vs Purchase Amount')
```



```
[191]: plt.figure(figsize=(4, 3))
sns.countplot(data=data, x = "Frequency of Purchases", hue="Gender")
plt.xticks(rotation = 45)
plt.show()
plt.tight_layout()
```



<Figure size 640x480 with 0 Axes>

Applying Techniques (t-test & greedy forward feature selector)

I want to get a deeper look into the difference between the genders and the frequency of purchases, so I will conduct a **t-test**.

For the **null hypothesis**: There is no significant difference between the two genders and their shopping frequency.

Alternative Hypothesis: There is a difference between the genders.

Before I begin, I need to **hot-encode** the Frequency variable, which will help me in developing future models as well!

The scale will be 1: Least Frequent (Annually) - 7: Most Frequent (Fortnightly)

```
[192]: #hot encoding the frequency variable.
data['Frequency Encoded'] = data['Frequency of Purchases'].map({
    'Annually': 1,
    'Quarterly': 2,
    'Every 3 Months': 3,
    'Monthly': 4,
    'Weekly': 5,
    'Biweekly': 6,
    'Fortnightly': 7
})
#displaying just to make sure it encoded properly
```

```
data[['Customer ID', 'Frequency of Purchases', 'Frequency Encoded']][1:4]
```

```
[192]:   Customer ID Frequency of Purchases  Frequency Encoded
1                  2             Fortnightly        7.0
2                  3              Weekly          5.0
3                  4              Weekly          5.0
```

```
[193]: #we will inject the median into the null values, since we now are able to use
       ↪the column numerically
data['Frequency Encoded'] = data['Frequency Encoded'].fillna(data['Frequency
       ↪Encoded'].median())
data['Frequency Encoded'].isnull().sum()
```

```
[193]: np.int64(0)
```

```
[194]: #now setting up the t-test:
females = data[data['Gender'] == 'Female']['Frequency Encoded']
males = data[data['Gender'] == 'Male']['Frequency Encoded']
ttest_ind(a=males,
          b=females,
          equal_var=False, alternative='two-sided')
```

```
[194]: TtestResult(statistic=np.float64(0.518038647587995),
pvalue=np.float64(0.6044772294810729), df=np.float64(2495.3900834759243))
```

My statistic number was about .53, meaning that females tend to buy *less* frequently than males.

The p-value was about .60, which is much higher than .05. This number being larger tells us that **there is no statistical significant difference** between the mean frequency of purchases between the two genders. We can **accept the null hypothesis**, and reject the alternative.

On the surface, it seemed like males purchased with more frequency, but using the t-test allowed us to see that there is in fact no statistical significant difference between the frequency of purchases between the genders.

We could keep searching for key features manually using tests like these, but a **greedy algorithm** could **save us some time**. It is very interesting to apply our in class techniques to a data set, and see how we can explore different relationships beyond the surface level.

```
[195]: # I will first encode the other features I want to use that aren't previously
       ↪numeric
data['Discount Encoded'] = data['Discount Applied'].map({
    'Yes': 1,
    'No': 0
})
data['Season Encoded'] = data['Season'].map({
    'Winter': 0,
    'Spring': 1,
```

```

        'Summer': 2,
        'Fall': 3
    })
data[['Discount Applied', 'Discount Encoded', 'Season', 'Season Encoded']][1675:
    ↪1678]

```

```
[195]:      Discount Applied  Discount Encoded  Season  Season Encoded
1675            Yes                  1  Spring           1
1676            Yes                  1  Winter           0
1677           No                   0  Summer           2
```

```
[196]: #setting up the greedy algorithm model to see what affects frequency of purchases
featured_cols = data[['Age', 'Purchase Amount (USD)', 'Previous Purchases',
    ↪'Review Rating', 'Discount Encoded', 'Season Encoded', ]]
X = featured_cols.values
y = data[['Frequency Encoded']].values.reshape(-1, 1)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)

# Define model
model = LinearRegression()

# Forward selection
sfs = SequentialFeatureSelector(model, n_features_to_select=3,
    ↪direction='forward')
sfs.fit(X_train, y_train)

# Selected features
print("Selected features:", sfs.get_support())
print(featured_cols.columns[sfs.get_support()])
```

```
Selected features: [ True False  True False False  True]
Index(['Age', 'Previous Purchases', 'Season Encoded'], dtype='object')
```

Based on the Sequential Feature Selector algorithm, from the set of columns I selected, we can see that the most influential factors for Frequency of Purchases are **Purchase Amount**, **Previous Purchases**, and **Review Rating**.

I selected the 6 columns based on eliminating some that I already looked at, and choosing ones that already had numerical values or were simple to encode.

With this information, I can now create a prediction model to predict how frequently a customer may make purchase, and I could redo this to find the most influential variables for the Purchase Amount.

Git Repo Link: <https://github.com/dpineda04/HiddenGems.git>